

Refactor or Rebuild

Whether AI Job Agent overbuilt company-level discovery, whether the backend should be refactored or restarted, and the simplest architecture that actually gets fresh jobs flowing — evaluated against 10 n8n templates and the project's own verified evidence.

Base evidence **The Ingestion Gap + The Ingestion Verdict, 2026-09-26**

New research **10 n8n templates + Bayt.com/JSearch/RSS providers, sourced**

Scope **decision only — nothing modified, activated, or committed**

Contents

01	Legend
02	The big question, answered
03	Confirmed current state
04	Template & provider research
05	Component-by-component verdict
06	Refactor vs. rebuild
07	The layered acquisition strategy
08	Can we stop company-level discovery?
09	Job quality & trust tiers
10	Matching architecture
11	Location normalization
12	Student vs. Pro eligibility
13	Cost model
14	Reliability & failure handling
15	Compliance & source risk
16	Two architectures, side by side
17	P0 / P1 / P2
18	Sequence to MVP
19	Founder decisions
20	Final verdict

· Legend

FACT verified in our code/DB/live n8n instance, or sourced external research

INFERENCE a reasoned conclusion built on facts, not independently verified

RECOMMENDATION what I'd do — actionable, but yours to override

FOUNDER DECISION requires your explicit sign-off before anything proceeds

01 The big question, answered directly

Did we overbuild the wrong layer? **Partially, yes** — and the evidence for it got stronger this pass, not weaker.

12.6%

REGISTRY COVERAGE
CEILING WITH THE CURRENT
COMPANY-FIRST, ADAPTER-
BUILT APPROACH

5/5

TARGET MARKETS A SINGLE
NEW PROVIDER (BAYT.COM)
COVERS NATIVELY, WITH
ZERO COMPANY RESEARCH

\$0

COST OF AN APIFY CAREER-
PAGE ATTEMPT THAT
DOESN'T MATCH A
SUPPORTED ATS

We spent real, substantial effort on **exhaustive upfront company-level discovery** — 588 CSV rows of hand-researched companies across five markets, a Registry Sync pipeline to keep growing that list, and a Source Intelligence Analyzer built specifically to figure out which ATS each of those 593 companies uses. That effort produced a registry that is genuinely well-built (real dedup, real idempotency, RLS-sound schema) — but it answers the wrong first question. It answers *"which companies exist and where's their careers page,"* when the business-critical question was always *"how do we get a large volume of fresh, trustworthy jobs flowing, starting now."*

Two new findings this pass make the "wrong layer" conclusion concrete rather than theoretical:

- **FACT** The Apify career-page actors identified in "The Ingestion Verdict" charge **nothing** for an attempt against an unsupported platform — meaning Source Intelligence Analyzer's core justification (avoid wasting paid attempts on sources that won't match) mostly evaporates. You don't need to classify a source before trying it; trying it *is* the free classification.
- **FACT** **Bayt.com** — the dominant native Gulf/MENA job board — has multiple independent, competitively-priced Apify actors (~\$0.75–0.80/1,000 results) covering all five of our target markets directly, with richer structured fields than any general aggregator researched. This didn't surface in the previous research pass and is the single most significant new fact in this report: **a huge share of "enough fresh jobs" may be answerable with zero company-level research at all.**

None of this means the registry work was wasted (Section 4 goes component-by-component) — it means the registry should never have been the *gate* ingestion waited behind. It belongs underneath a broader, provider-first acquisition layer as a trust/provenance mechanism, not in front of it as a prerequisite.

02 Confirmed current state

Carried forward from "The Ingestion Gap" and "The Ingestion Verdict" (same session). Re-confirmed via `git status` this pass: **zero changes to the repository since those audits** — same branch, same HEAD, same uncommitted working tree. Nothing has drifted; the earlier evidence is still current.

- **Registry:** 559 companies, 593 company_sources. 17/593 (2.9%) ingestible by existing in-house adapters (Greenhouse/Lever/Workable). 75/593 (12.6%) on any verified ATS. Zero rows for the International Remote market.
- **Ingestion:** Job Ingestion — Pilot Orchestrator exists, well-tested, hardcoded to 4 sources, never run outside a local pilot. Zero jobs in production.
- **n8n:** 5 workflows in repo, 15 on the live shared instance. CV Analysis Worker has drifted (live "v2" vs. committed older version). Three undocumented shadow workflows exploring AI scraping and LinkedIn-via-Apify were created the same day as the first audit, outside version control.
- **Downstream:** matching, cover-letter generation, and application sending remain entirely unbuilt, despite a ready RLS-sound schema for all three since 2026-08-09.
- **Frontend/auth/CV analysis:** not re-audited this pass — explicitly out of scope for this architecture question, per your own framing. Nothing here suggests touching them.

03 Template & provider research

All 10 requested templates researched, plus a broader sweep for adjacent providers.

TEMPLATE / PROVIDER	WHAT IT ACTUALLY IS	IMPACT
1. Careers-page CSV leads (Apify)	Domain-in, jobs-out via an undisclosed Apify actor. Confirms the pattern, adds no new capability beyond dami_studio/pulsedata already known.	LOW
2. Snapshot Site + GPT-4o career extraction	Snapshot Site is a real JS-rendering node (not Apify) — genuinely reaches the JS-heavy custom-HTML bucket career-page actors can't. GPT-4o structures the output with zero schema validation as published — a real hallucination-risk gap, not a solved problem.	MEDIUM-HIGH
3. JSearch (RapidAPI)	Aggregates via Google for Jobs → LinkedIn, Indeed, ZipRecruiter, Glassdoor, and Bayt.com . \$25/mo for 10k requests. Location is country + free-text city only (no confirmed remote flag); apply-URL canonicity unconfirmed.	HIGH, with caveats
4. AI-ranked digest, 6+ boards	Real sources: RemoteOK, Remotive, Arbeitnow, Jobicy, Himalayas, TheMuse. Demo-level: no persistence, no schema-validated AI output, no eligibility pre-filter. The shape is right; none of the hardening is present.	MEDIUM
5. Indeed + Tavily	Actually a B2B sales-lead tool filtered to sub-250-employee companies — wrong product shape entirely for a jobseeker platform.	LOW
6. LinkedIn Jobs (Apify, no cookies)	42 structured fields incl. pre-parsed location — richest data of anything researched. No login required. Not a recommendation to use — see Section 18.	Separate decision
7. RSS remote-job digest	RemoteOK, WeWorkRemotely, Himalayas confirmed. A related aggregator explicitly documents "fields a feed doesn't publish are null, never guessed" — independently validates our own location-integrity rule.	MEDIUM-HIGH for Intl. Remote
8–9. Upwork templates	Upwork itself irrelevant. The pipeline shape (trigger → Apify → normalize → dedup → store → deliver) recurs identically across nearly every template studied — strong independent confirmation this is the right shape.	LOW (pattern only)
10a. Multi-ATS actors (other vendors)	Independent actors confirm the same 9–12 platform list via official public APIs — reduces single-vendor risk. One vendor discloses Workable's public endpoint returns zero jobs for every account they tested .	MEDIUM (risk note)
10b. Bayt.com	The dominant native Gulf/MENA job board. Multiple independent Apify actors, 41–65+ fields, 15–18 MENA markets, ~\$0.75–0.80/1,000 results. Covers all 5 target markets directly — first-party regional data, not secondhand aggregation.	HIGH — standout finding
10c. Adzuna	Confirmed 18–19 country coverage — zero Middle East presence. Rule out entirely for this product.	None — dead end

04 Component-by-component verdict

No component is deprecated without evidence. No component survives untouched just because it already exists.

COMPONENT	VERDICT	WHY
companies	KEEP AS-IS	Canonical-identity concept is needed regardless of ingestion model.
company_sources	KEEP, CHANGE ROLE	From mandatory gate to trust/provenance layer, populated reactively rather than exhaustively upfront. The 593 existing rows become the seed of that trust layer.
jobs	KEEP, EXTEND	Schema is sound (nullable <code>source_id</code> already tolerates provider-sourced jobs with uncertain provenance). Needs additive columns: <code>trust_tier</code> , <code>provider</code> , <code>canonical_apply_url</code> quality, structured location fields.
matches, cover_letters, applications, notifications	KEEP AS-IS	Schema exists, RLS-reviewed, zero code built. They just need their workers built.
automation_tasks	KEEP, SIMPLIFY AMBITION	No claim function built yet. Don't generalize further before real volume justifies a queue abstraction.
Registry Sync	KEEP, CHANGE ROLE	From "grow the registry that gates ingestion" to "add a trusted first-party source, proactively or reactively." Its 3-adapter pattern stays useful unchanged.
Source Intelligence Analyzer	KEEP, REDUCE FURTHER	Since unsupported-platform Apify attempts are free, there's no cost reason to classify before trying. Demote to a low-priority, opportunistic tagger, off the critical path.
Job Ingestion — Pilot Orchestrator	KEEP, CHANGE ROLE	Becomes specifically the Tier A runner. Its retry/backoff/dedup machinery is the reference implementation every new tier's workflow should copy. Activate now.
ATS adapters (Greenhouse/Lever/Workable)	KEEP the 3 that exist; REPLACE the plan to build more	Zero marginal cost forever. Stop building new in-house adapters — Apify's 9–12-platform coverage makes that a poor trade now.
Error Handler	KEEP AS-IS	Unaffected by this pivot. Still P1: wire the notification boundary.
CSV company files	KEEP	Seed data, not obsolete. Unimported International Remote rows remain real, unconsumed value.
source_intelligence tables/migrations	KEEP, FREEZE FURTHER INVESTMENT	Real, tested infrastructure — not ripped out. Stop expanding its promotion/retry logic; it's less load-bearing than originally planned.

05 Refactor vs. rebuild — a direct recommendation

DIMENSION	OPTION 1 — REFACTOR	OPTION 2 — CLEAN RESTART
Engineering time	Additive schema changes + new ingestion tiers. Auth/CV Worker/frontend untouched.	Re-implement auth integration, RLS, dedup, idempotency — already correct today.
Migration risk	Low — additive, reviewed, reversible.	High — a second schema must interoperate with the kept frontend/CV Worker.
Data loss risk	None — registry becomes the trust layer.	Real — 593 researched companies get ported or abandoned.
Complexity introduced	One new router concept + provider integrations, added to a working system.	A second backend running alongside/replacing the first.
Technical debt	Reduced — the debt being fixed is the registry-as-gate assumption.	New debt nearly guaranteed under time pressure.
Testing burden	Existing dedup/retry/idempotency suites remain valid.	All existing ingestion tests void; suite rebuilt from zero.
Time to MVP	Fast — Pilot Orchestrator ships this week regardless.	Slow — nothing ships until parity is reached.
Existing schema quality	Real idempotency, real dedup fix, RLS-sound, already tolerates non-registry-sourced jobs.	Would have to be re-proven from scratch.
Code reusability	High — retry pattern, dedup logic, RLS, Error Handler, CV Worker all reusable.	Low by construction.
Rollback ability	High — each tier independently addable/removable.	Low — all-or-nothing cutover.

Recommendation: Option 1 (Refactor), not Option 2

Every dimension favors refactor, and no technical evidence forces a rebuild — the schema already tolerates the model this report recommends. Mapping to the original A/B/C/D framing: the right answer is **closest to A combined with a narrow, evidence-scoped version of B** — keep the database/backend foundation and simplify/re-role the automation layer (A), while replacing the *ingestion mechanism* around a layered provider strategy rather than continued in-house build-out (a targeted slice of B). **Reject C** (clean restart) — no evidence justifies the cost. **Reject D alone** — the ingestion model's core assumption needs to flip, which is a real re-architecture of one layer, not incremental patches.

06 The layered acquisition strategy

Yes, this is better than `company_registry` → `source_intelligence` → `adapters` → `jobs` — and `company_sources` still sits underneath it, as provenance/control, not as a gate.

TIER	WHAT IT IS	ROLE GOING FORWARD
A	In-house ATS adapters (Greenhouse, Lever, Workable)	Keep, zero marginal cost. Don't extend with more in-house builds.
B	Apify multi-ATS career-page extraction (9–12 platforms, free on no-match)	Primary path for any company on a known-but-unadapted ATS. Replaces the "build more adapters" plan.
C1	Bayt.com via Apify — native MENA job board, all 5 markets	New, primary broad-supply layer. Test first — closest regional fit of anything researched.
C2	JSearch (RapidAPI) — secondary broad aggregator, confirmed MENA reach via Bayt	Complementary breadth layer once Bayt is validated. Verify apply-URL canonicity first.
D	RSS/free remote boards (RemoteOK, WeWorkRemotely, Himalayas, Remotive, Arbeitnow, Jobicy, TheMuse)	Closes the confirmed zero-coverage International Remote gap, near-zero cost.
E	AI/browser fallback (Snapshot Site + GPT-4o pattern, schema-validated)	Last resort for JS-heavy custom pages with no known ATS. Discovery-only, always schema-validated.
F	Manual / not automatable	Unchanged.

`company_sources` stops being "the thing every job must trace back to before it can be ingested" and becomes "the record of which employers we trust enough to treat as canonical/first-party." A job arriving via Tier C1/C2/D can exist and be shown to users at a lower trust tier without a `company_sources` row; a row only gets created to upgrade that employer's trust level or guarantee freshness.

07 Can we stop company-level discovery?

Mostly yes, for volume. Not entirely, for trust.

INFERENCE The evidence supports flipping the primary flow to **Sources/Providers** → **raw jobs** → **normalize** → **company canonicalization (after ingestion)** → **trust-tier tag** → **dedupe** → **store**, rather than **Companies** → **find source** → **find ATS** → **scrape**. Canonicalization-after-ingestion means: fuzzy-match the employer name/domain against `companies`, create a new canonical row if no confident match exists (reusing Registry Sync's existing `resolve_registry_candidate()` pattern), and only escalate to a full `company_sources` row when that employer is worth first-party treatment.

What becomes **optional**: exhaustive upfront research to enumerate every possibly-relevant company before ingestion can start. That work no longer gates volume.

What **still requires** the company registry: guaranteeing freshness/trust for specific high-value employers, preventing duplicate/conflicting employer identities across tiers, and providing the canonical-priority tie-break when the same job arrives via two different tiers. The registry doesn't disappear — it moves from a blocking prerequisite to an enrichment layer built out where it earns its cost.

08 Job quality & trust tiers

TRUST TIER	EXAMPLES	DIRECT WRITE TO CANONICAL JOBS?
1. First-party ATS	In-house Greenhouse/Lever/Workable; Apify-detected supported-ATS extraction	Yes — highest trust
2. First-party careers page, verified	Custom HTML, human or schema-validated AI extraction	Yes, after validation
3. Reputable provider, canonical URL confirmed	Bayt.com, JSearch (once verified)	Yes, tagged with provider
4. Public job board, no guaranteed canonical URL	Generic aggregator with only a tracking/redirect link	Discovery-only until resolved
5. AI/browser discovered	Snapshot Site + GPT-4o pattern, any Tier E output	Never without the same deterministic validation every other tier passes

Canonical apply URL: first-party employer domain > verified ATS URL > provider apply-link resolving to the employer > never a bare aggregator search-result link.

Employer identity: fuzzy-match name+domain against `companies` ; no confident match holds the candidate for review, never silent auto-create.

Stale/fake/duplicate rejection: require a freshness signal where available; cross-tier dedup on a normalized (company, title, location) fingerprint; low-trust jobs with no posting date get a shorter default expiry.

09 Matching architecture

Simpler than planned, per the evidence — but not as simple as the templates make it look.

Template #4 (6+ boards, Anthropic digest) validates the core shape: normalize-many-sources → filter → LLM scoring with a structured contract (score, reason, missing skills/risk, tailored pitch) → deliver. Directly reusable as a starting schema.

What it does **not** validate, and what AI Job Agent still needs regardless of this pivot (AGENTS.md §30):

- **Hard eligibility pre-filter before any LLM call** — never an LLM's decision, also the cost control.
- **Schema-validated, tool-call-constrained AI output** — not a plain prompt hoping for structured text.
- **Persisted match records with reasoning stored** — not an ephemeral one-shot digest.
- **Idempotency** — the same job×user pair must not be rescored repeatedly.
- **Explicit "not eligible, here's why"** handling, never silent omission.
- **Per-user/per-run cost caps.**

No researched template is production-ready on these points. The reusable part is the shape; the production-grade parts still have to be built.

10 Location normalization architecture

Every source's raw location must pass through one deterministic normalizer before a job is eligible for anything plan-related. Never guess.

QUESTION	ANSWER
Does each source return structured location?	FACT Inconsistent. JSearch: country + free-text city, no confirmed remote flag. RSS: explicitly null when unpublished. Bayt.com: richer structured fields. LinkedIn: pre-parsed (not in use).
Can it be deterministically normalized otherwise?	Yes, via a gazetteer mapping known cities (Dubai→AE, Riyadh→SA, Doha→QA, Kuwait City→KW, Beirut→LB) plus a controlled remote-scope vocabulary. Anything that doesn't match cleanly is flagged, not guessed.
What about "MENA" / "Worldwide" / bare "Remote"?	Map to a controlled <code>remote_scope</code> (country / regional_mena / worldwide), distinct from <code>country_code</code> . A bare "Remote" with no scope is held for review, not assumed worldwide.
Storage?	<code>country_code</code> (ISO, nullable), <code>city</code> (free text, nullable), <code>work_arrangement</code> (onsite/hybrid/remote), <code>remote_scope</code> , <code>location_confidence</code> (high/low).
Preventing bad eligibility from bad parsing?	Eligibility queries filter on <code>location_confidence='high'</code> only. Low-confidence/null-country jobs are excluded until resolved — fail closed, never soft-warn.

11 Student vs. Pro eligibility

PLAN	ELIGIBILITY RULE
Student	<code>country_code='LB'</code> AND onsite/hybrid, OR remote AND <code>remote_scope='country'</code> AND <code>country_code='LB'</code> . A bare "Remote" with no confirmed Lebanon scope is excluded.
Pro	Lebanon + Gulf/MENA (5 target markets) + international remote (<code>remote_scope</code> regional_mena/worldwide) + relocation-aware only via explicit user opt-in, never inferred from the job.

Enforced server-side, never client-side; fails closed on any `location_confidence='low'` job regardless of plan.

12 Cost model

TIER	REAL COST, SOURCED
A — in-house adapters	\$0
B — Apify multi-ATS	\$0.55–\$1 / 1,000 extracted jobs; \$0 on no-match
C1 — Bayt.com	~\$0.75–\$0.80 / 1,000 results
C2 — JSearch	\$25/mo for 10,000 requests; 200 free/month
D — RSS	\$0
E — AI/browser fallback	Rendering + GPT-4o token cost, not independently priced — use sparingly by design

INFERENCE With B, C1, and D active, reaching genuinely useful volume across all 5 markets plus international remote is very plausibly a **low-double-digit-dollars-per-month problem**, not the mid-hundreds implied by continuing in-house adapter builds. This materially de-risks the "enough jobs" business risk.

13 Reliability & failure handling

Unchanged principles from "The Ingestion Verdict," extended to the new tiers: per-source/per-provider health tracking, never guess on a failed fetch, distinguish "failed" from "succeeded empty," bounded retries, auto-pause a repeatedly-failing source.

- **FACT** An independent Apify vendor discloses **Workable's public widget endpoint returned zero jobs for every account they tested**. Since Workable is one of our three in-house adapters, this is a concrete, cheap, worth-doing check on our own Workable pilot source, independent of anything else in this report.
- JSearch and Bayt.com both need their own rate-limit/quota tracking — the same `last_provider_run_at` / `provider_cost_class` fields proposed in "The Ingestion Verdict" extend cleanly to cover them.

14 Compliance & source risk

SOURCE	RISK POSTURE
Apify career-page/ATS extraction	Low — public company career pages, no authentication bypass.
Bayt.com	Low-moderate — a public commercial job board scraped by multiple independent vendors openly for years; standard ToS caveats apply, no special red flag found.
JSearch	Low — licensed commercial API. Worth a founder-visible footnote: it aggregates via Google for Jobs, which includes LinkedIn among its underlying sources. Not us scraping LinkedIn, but not zero-nexus either.
RSS feeds	None — publisher-syndicated by design.
AI/browser fallback (Tier E)	Standard hallucination/validation risk, already governed by AGENTS.md §30 — not a new risk class.
LinkedIn direct scraping	Unchanged — explicitly not enabled, remains a separate founder/compliance decision.

15 Two architectures, side by side

Current

```
CSV Research (588 rows, 5 markets)
  ↓
Registry Sync (manual / Sheets / Apify-dataset-read)
  ↓
company_sources (593 rows) ← GATE: nothing downstream runs without a row here
  ↓
Source Intelligence Analyzer (free GET, classify ats_provider)
  ↓ (0 high-confidence hits across 60 real observations)
Manual: research + build a new in-house ATS adapter, one platform at a time
  ↓
Job Ingestion – Pilot Orchestrator (3 adapters, 4 hardcoded sources)
  ↓
jobs (20 test rows, 0 in production)
  ↓
??? – matching / cover letters / applications: UNBUILT
```

Recommended simplified

```
Tier A In-house adapters (Greenhouse/Lever/Workable)      |
Tier B Apify multi-ATS extraction (9-12 platforms)         |
Tier C1 Bayt.com (native MENA, all 5 markets)              |— converge
Tier C2 JSearch (broad secondary aggregator)               |
Tier D RSS / free remote boards (Intl. Remote)            |
Tier E AI/browser fallback (schema-validated, discovery-only)|
  ↓
Normalize → Location-normalize (deterministic, null if unclear) → Trust-tier tag
  ↓
Company canonicalize (fuzzy-match / create) → Deduplicate → Canonical source resolution
  ↓
jobs (trust_tier, provider, canonical_apply_url, country_code, work_arrangement, remote_scope)
  ↓
Source Health Rollup →→ feeds back into company_sources
  (now: trust/provenance layer, populated reactively – not a gate)
  ↓
Matching Worker (hard eligibility filter first → schema-validated score+reason → persisted)
  ↓
Notifications / Cover Letters / Applications
```

WHAT DISAPPEARS BETWEEN THEM

- Exhaustive upfront company-level research as a **prerequisite** — the CSVs and 593 rows survive, but nothing waits on growing that list further.
- Source Intelligence Analyzer as a **required** pre-ingestion step — demoted to optional, off the critical path.
- The plan to build additional in-house ATS adapters — replaced by Apify's broader platform coverage.
- `company_sources` as a **blocking dependency** — survives entirely, but as a trust layer underneath the flow.

16 P0 / P1 / P2

PRIORITY	ITEMS
P0	• Activate Job Ingestion Pilot Orchestrator (17 sources, scheduled) — free value on the table • Reconcile CV Analysis Worker repo/live drift; freeze the shadow n8n workflows • Run the Apify career-page mini-benchmark already planned • New: pilot Bayt.com + JSearch trial + a small RSS pull in parallel — cheap, fast, highest-leverage new signal in this report
P1	• Begin the minimal matching worker against whatever real jobs exist from any tier so far • Design (not migrate) the trust-tier + location-normalization schema additions • Workable adapter health-check (cheap, fast, concrete disclosed risk) • Error Handler wiring; source health tracking extended to new provider tiers • International Remote coverage via RSS (near-zero cost, closes a confirmed gap)
P2	• Any further in-house ATS adapter — likely unnecessary given Apify's coverage • Cross-tier/cross-source dedup — only matters once ≥ 2 tiers are live • Cover-letter and application-sending flows — after matching and a wider job set are live • LinkedIn/Apify founder decision — parked

17 Sequence to MVP

1. Activate Pilot Orchestrator (Tier A, 17 sources, scheduled, Error Handler attached).
2. Reconcile CV Analysis Worker drift and the shadow workflows — governance, not new capability.
3. Run the Apify career-page benchmark and, in parallel, a small Bayt.com + JSearch + RSS pilot — both cheap, both read-only until results are in.
4. Build the minimal matching worker against whatever real jobs exist at this point, from any tier — do not wait for the full provider mix to be finalized.
5. Using combined benchmark results, design the additive schema migration: trust_tier, provider, location fields, company_sources role change.
6. Apply the migration (reviewed, reversible) and wire the validated tiers into a real scheduled ingestion pipeline.
7. Decide on any remaining in-house adapter work only for gaps the full tier mix still can't cover.
8. Begin cover-letter and application-sending flows once matching and a genuinely multi-tier job set are both live.
9. Resolve the LinkedIn/Apify founder decision on its own track — not blocking anything above.

18 Founder decisions required

COMPLIANCE

LinkedIn-via-Apify remains unchanged from "The Ingestion Verdict" — richer data than anything else researched, but requires an explicit AGENTS.md §7 amendment plus real legal review before any use.

PROVIDER SPEND

Approve a small test budget to pilot Bayt.com and JSearch alongside the already-planned Apify benchmark — all three are cheap (low tens of dollars combined) but all three spend real money.

SCOPE

Confirm the registry's new role (reactive trust layer, not upfront gate) matches your product intuition.

PRIORITY

Matching worker vs. continued ingestion breadth — this report leans toward starting the matching worker immediately in parallel, not sequentially after ingestion is "done."

19 Final verdict

Refactor, not rebuild. Simplify the ingestion layer's role, don't discard its engineering.

The company registry, Source Intelligence Analyzer, and adapter-building plan were solving a real problem with real engineering discipline — they just weren't solving the problem that actually determines whether this business works. That problem is volume and freshness, and it turns out to be answerable far more cheaply than the current architecture assumed, largely because of one fact that wasn't visible until this pass: Bayt.com already exists, already covers all five target markets, and already charges commodity prices for exactly the data this product needs.

None of that makes the last month of work wasted. The schema is sound. The dedup and retry logic is genuinely good and directly reusable. The registry becomes the trust layer a mature version of this product needs anyway — it just stops being the thing everything else has to wait behind. Ship the Pilot Orchestrator this week, benchmark Bayt/JSearch/RSS alongside it, and start the matching worker in parallel rather than after — that sequence gets to a real company→jobs→matching→user-value flow faster than either "keep building the registry" or "start over" would.

Built on the confirmed evidence in "The Ingestion Gap" and "The Ingestion Verdict" (2026-09-26, re-confirmed unchanged via `git status` this pass) plus new, sourced research into 10 n8n templates and the Bayt.com/JSearch/RSS provider landscape. No code, database, credentials, schedules, branches, or n8n workflows were modified, activated, created, or deleted in the production of this report.